

Big Data Step-by-Step

Boston Predictive Analytics
Big Data Workshop

Microsoft New England Research &
Development Center, Cambridge, MA

Saturday, March 10, 2012

by **Jeffrey Breen**

President and Co-Founder
Atmosphere Research Group
email: jeffrey@atmosgrp.com
Twitter: @JeffreyBreen



<http://atms.gr/bigdata0310>

Big Data Infrastructure

Part I: Local VM

Code & more on github:

<https://github.com/jeffreybreen/tutorial-201203-big-data>

Overview

- Download and install a virtual machine containing a configured and working version of Hadoop
- Install R
- Copy some data into the HDFS
- Test our installation by running some small Hadoop jobs
- Extra credit: install RStudio

Thank you, Cloudera

- Cloudera's Hadoop Demo VM provides everything you need to run small jobs in a virtual environment
- Hadoop 0.20 + Flume, HBase, Hive, Hue, Mahout, Oozie, Pig, Sqoop, Whirr, Zookeeper
- Based on CentOS 5.7 & available for VMware, KVM and VirtualBox:

<https://ccp.cloudera.com/display/SUPPORT/Cloudera%27s+Hadoop+Demo+VM>

- Older versions came with training exercises, but fortunately they're still available on github:

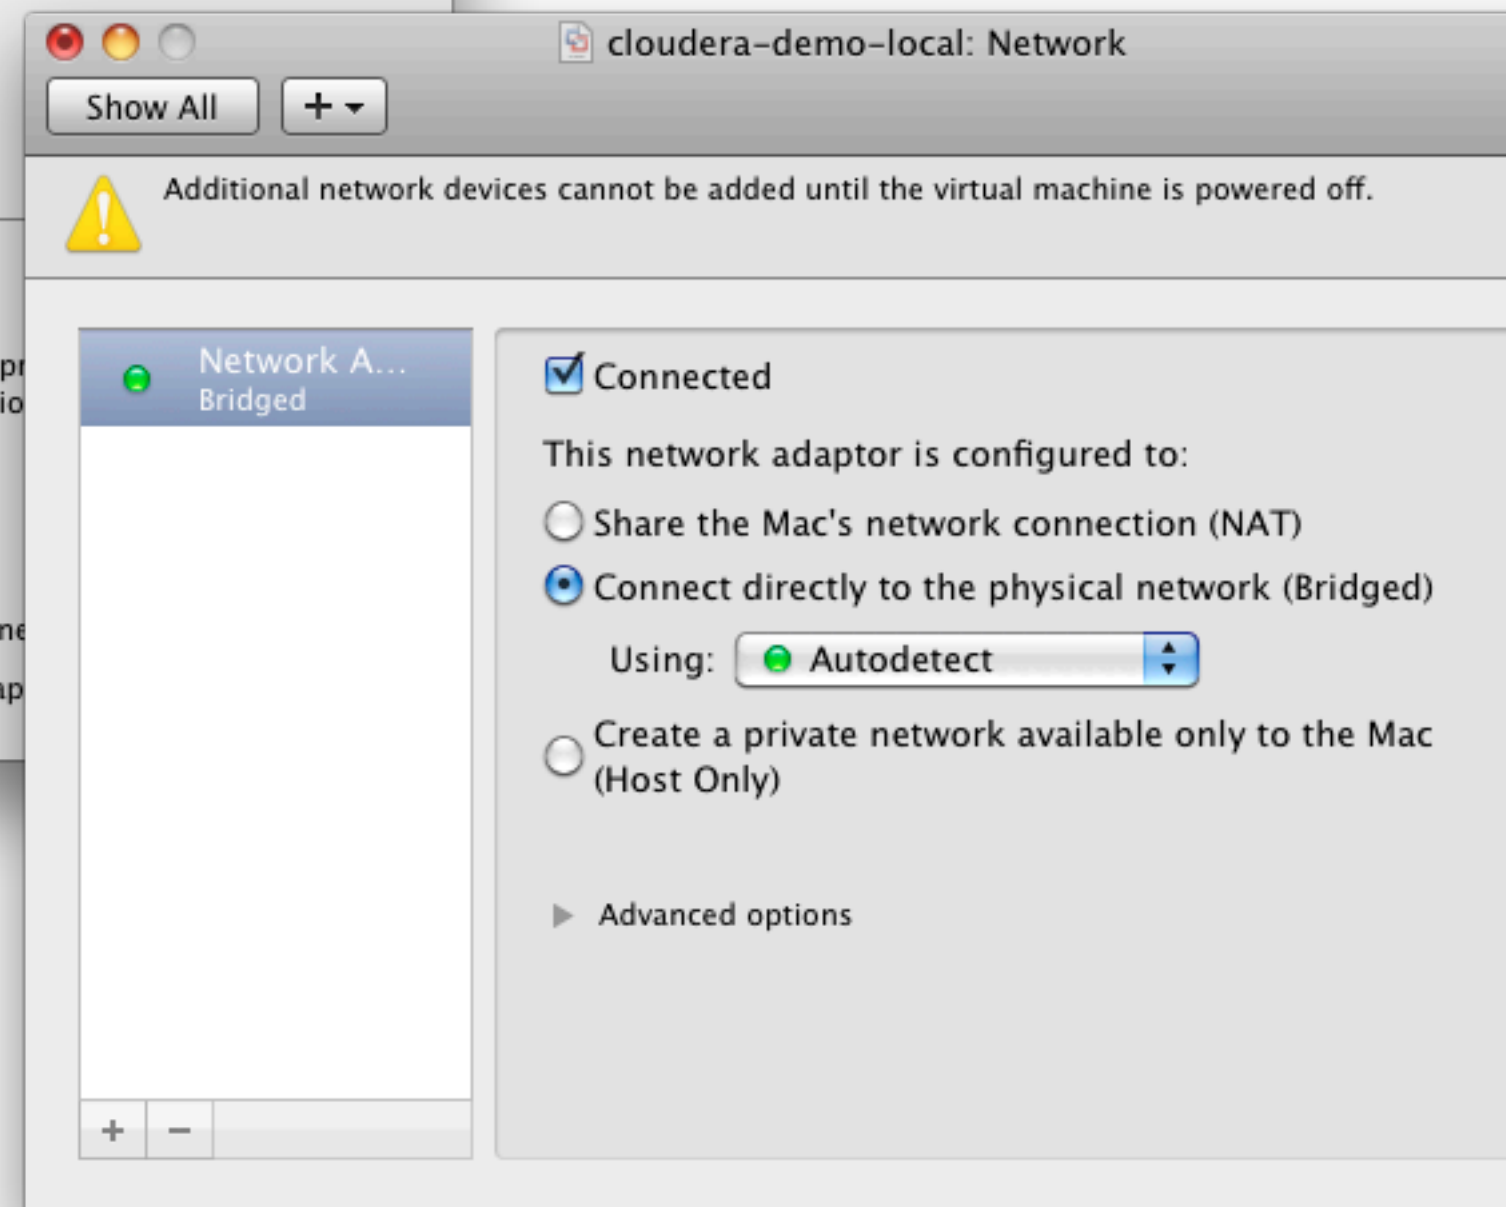
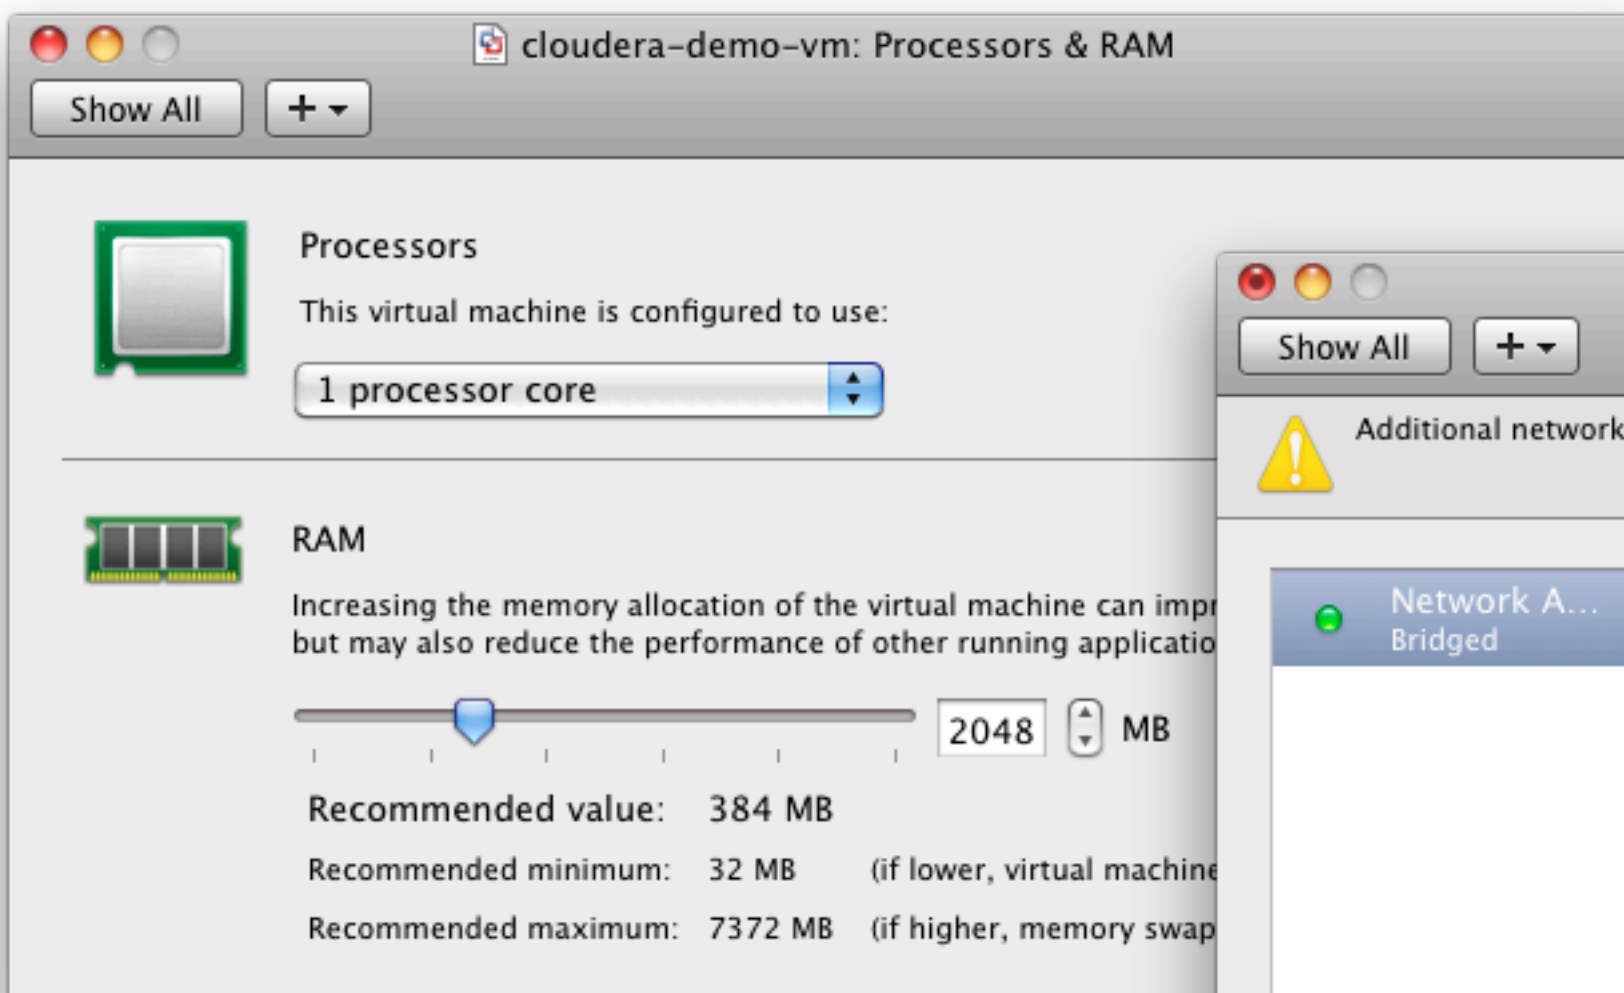
<https://github.com/cloudera/cloudera-training>

- Provides a common base which we will use for our later cluster, etc. work

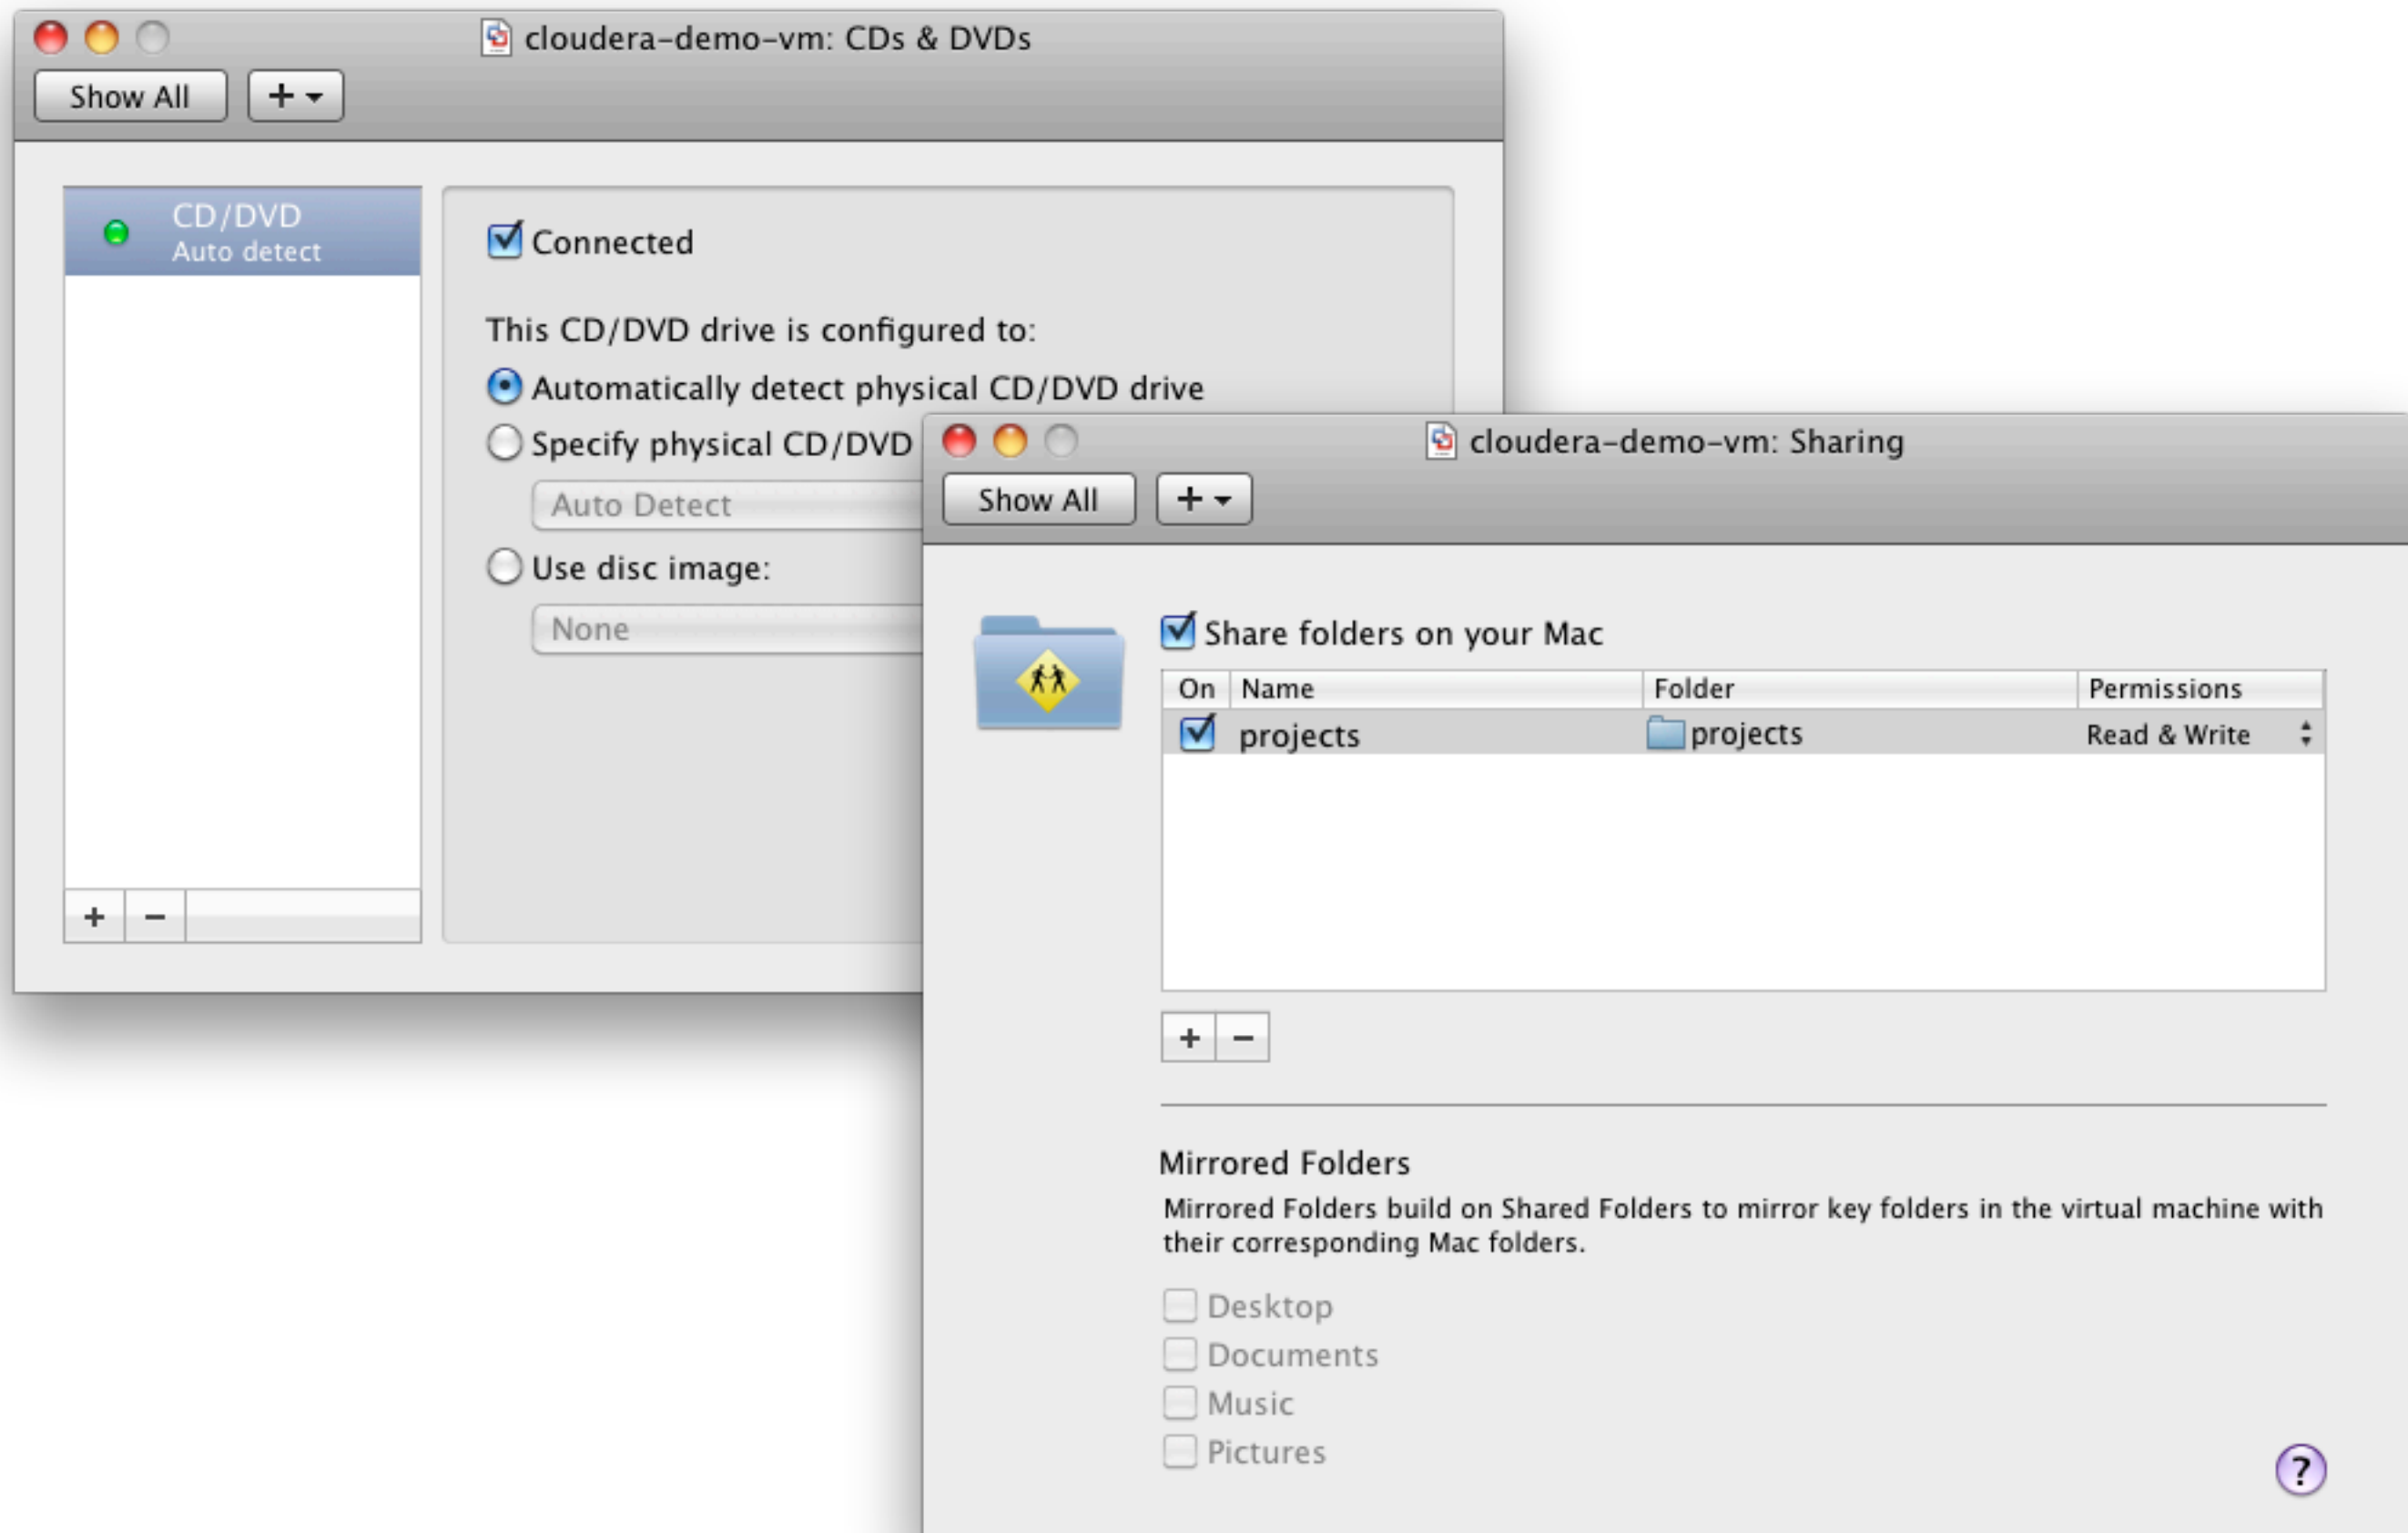
A couple of tweaks

- Give it more RAM
 - uses 1GB by default
 - not configured with a swap file
- Use Bridged networking vs. NAT or Host-only
 - Virtual machine will get its own IP address on your network
 - Experienced DNS errors with whirr while sharing an IP
- Extras: Set up shared folders & add a CD-ROM
 - Shared folders make it easy to share data & code between your computer and the VM
 - Add a CD-ROM drive if you want to install VMware tools or any ISO file

Important



Nice to have



Yes, it's that easy

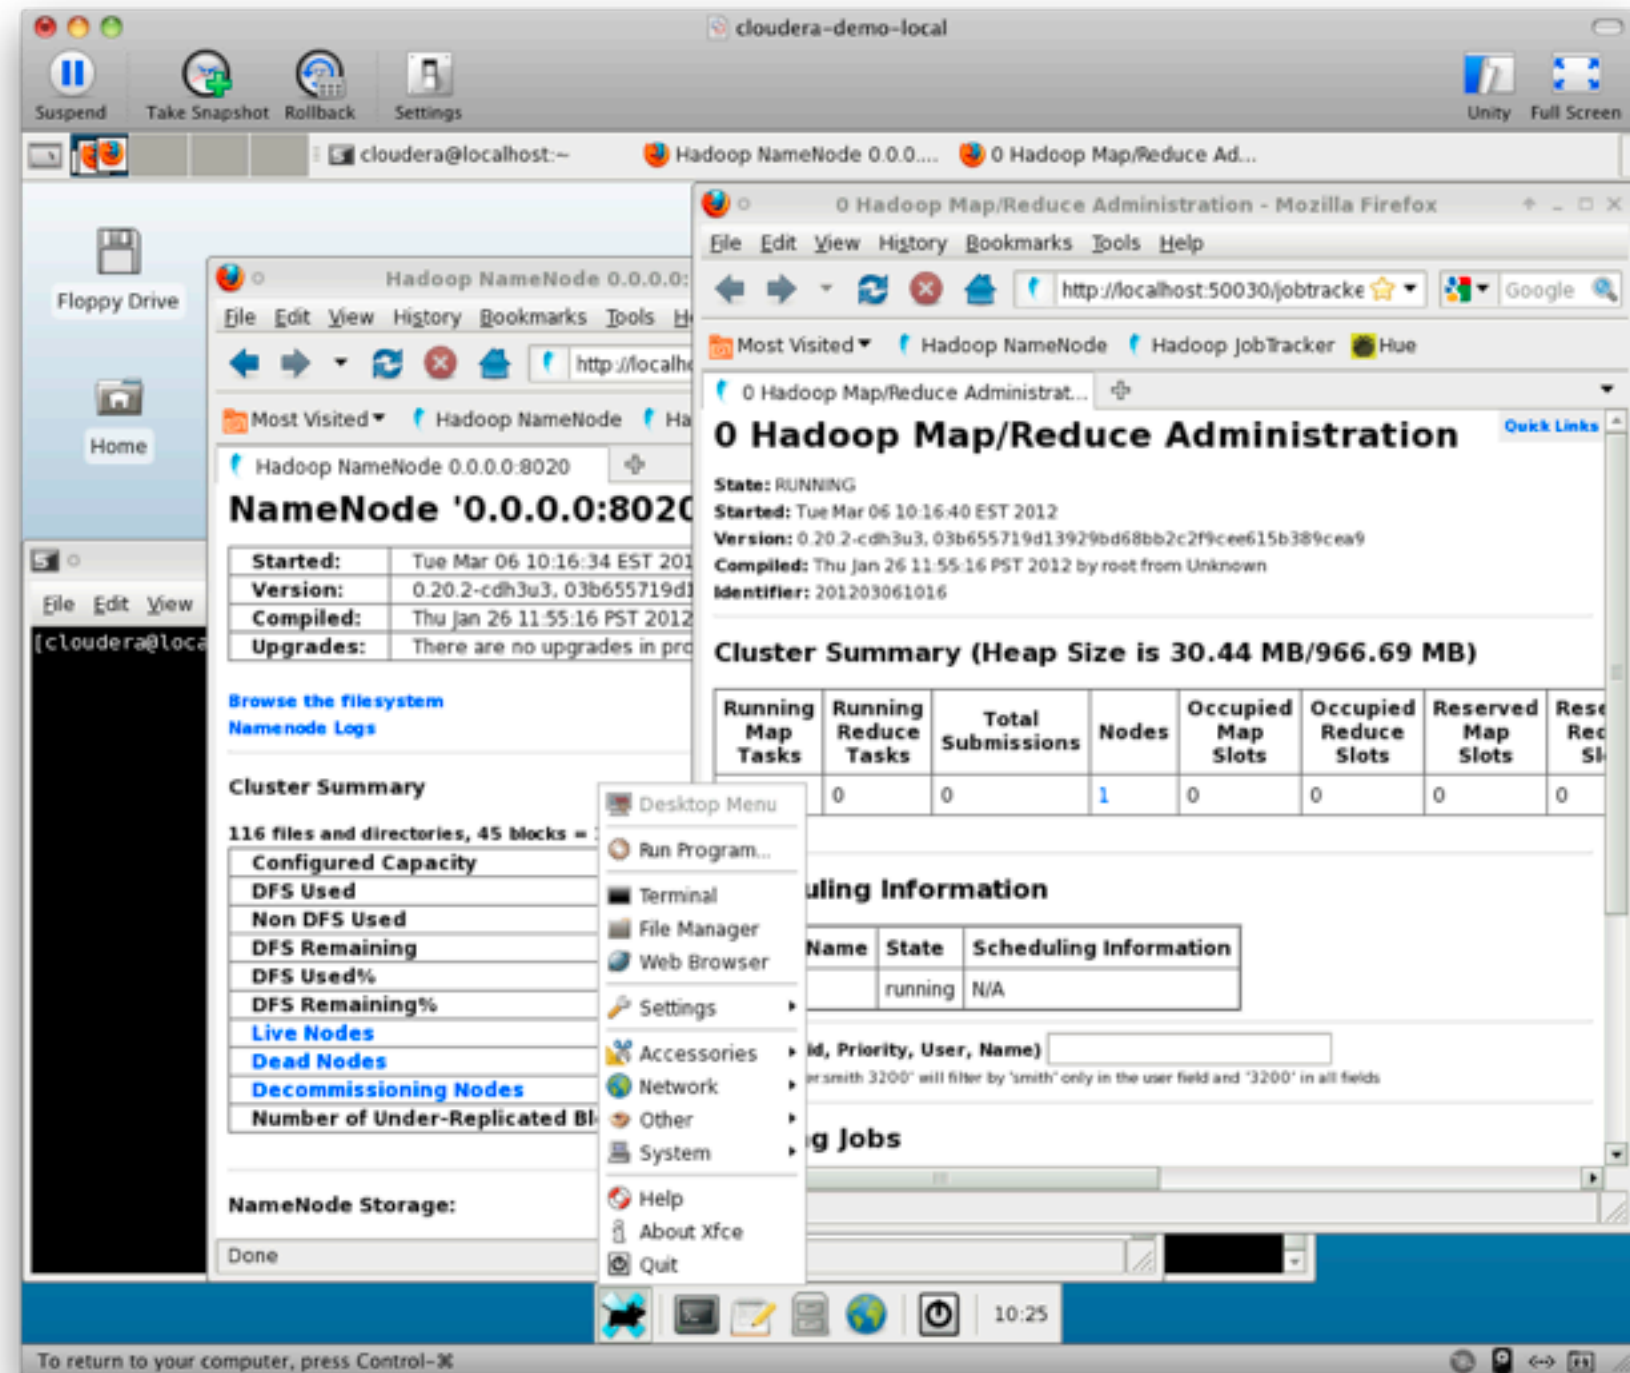
Boot VM and log in as “cloudera”. (Password = “cloudera” too)

Execute as root with “sudo”

“sudo su -” for root shell

Hadoop already running

Firefox contains bookmarks to admin pages



Well, almost.

- Install VMware tools and link to shared folder on host PC

```
$ sudo mkdir /mnt/vmware
$ sudo mount /dev/hda /mnt/vmware
$ tar zxf /mnt/vmware/VMwareTools-8.4.7-416484.tar.gz
$ cd vmware-tools-distrib/
$ sudo ./vmware-install.pl
$ ln -s /mnt/hgfs/projects/tutorial-201203-big-data/ ~/.
```

- Install handy utilities (wget, git)

```
$ sudo yum -y install wget git
```

- Install EPEL repository

```
$ sudo rpm -Uvh http://dl.fedoraproject.org/pub/epel/5/x86\_64/epel-release-5-4.noarch.rpm
```

- Install R from EPEL

```
$ sudo yum -y install R
```

- Set Hadoop environment variables (*workaround for CDH3u3 VM*)

```
$ sudo ln -s /etc/default/hadoop-0.20 /etc/profile.d/hadoop-0.20.sh
$ cat /etc/default/hadoop-0.20 | sed 's/export //g' > ~/.Renviron
```

Warning: Pages of fast-
scrolling gibberish to follow

But it's all going to be OK

```
[cloudera@localhost ~]$ sudo mkdir /mnt/vmware
[cloudera@localhost ~]$ sudo mount /dev/hda /mnt/vmware
mount: block device /dev/hda is write-protected, mounting read-only
[cloudera@localhost ~]$ tar xzf /mnt/vmware/VMwareTools-8.4.7-416484.tar.gz
[cloudera@localhost ~]$ cd vmware-tools-distrib/
[cloudera@localhost vmware-tools-distrib]$ sudo ./vmware-install.pl
Creating a new VMware Tools installer database using the tar4 format.

Installing VMware Tools.

In which directory do you want to install the binary files?
[/usr/bin]

What is the directory that contains the init directories (rc0.d/ to rc6.d/)?
[/etc/rc.d]

What is the directory that contains the init scripts?
[/etc/rc.d/init.d]

In which directory do you want to install the daemon files?
[/usr/sbin]

In which directory do you want to install the library files?
[/usr/lib/vmware-tools]

The path "/usr/lib/vmware-tools" does not exist currently. This program is
going to create it, including needed parent directories. Is this what you want?
[yes]

In which directory do you want to install the documentation files?
[/usr/share/doc/vmware-tools]

The path "/usr/share/doc/vmware-tools" does not exist currently. This program
is going to create it, including needed parent directories. Is this what you
want? [yes]

The installation of VMware Tools 8.4.7 build-416484 for Linux completed
successfully. You can decide to remove this software from your system at any
time by invoking the following command: "/usr/bin/vmware-uninstall-tools.pl".

Before running VMware Tools for the first time, you need to configure it by
invoking the following command: "/usr/bin/vmware-config-tools.pl". Do you want
this program to invoke the command for you now? [yes]

Initializing...

Making sure services for VMware Tools are stopped.

Stopping VMware Tools services in the virtual machine:
  Guest operating system daemon:      [ OK ]
  Virtual Printing daemon:           [ OK ]
  Unmounting HGFS shares:            [ OK ]
  Guest filesystem driver:            [ OK ]

Found a compatible pre-built module for vmxmemctl.  Installing it...

Found a compatible pre-built module for vmhgfs.  Installing it...
```

```
[cloudera@localhost ~]$ sudo yum -y install wget git
Loaded plugins: fastestmirror
Loading mirror speeds from cached hostfile
 * base: mirror.symnds.com
 * epel: mirror.symnds.com
 * extras: mirrors.einstein.yu.edu
 * updates: mirror.symnds.com
epel | 3.4 kB 00:00
epel/primary_db | 3.7 MB 00:01
Setting up Install Process
Resolving Dependencies
There are unfinished transactions remaining. You might consider running yum-complete-transaction first to finish them.
The program yum-complete-transaction is found in the yum-utils package.
--> Running transaction check
---> Package git.x86_64 0:1.7.4.1-1.el5 set to be updated
--> Processing Dependency: perl-Git = 1.7.4.1-1.el5 for package: git
--> Processing Dependency: perl(Error) for package: git
--> Processing Dependency: perl(Git) for package: git
---> Package wget.x86_64 0:1.11.4-2.el5_4.1 set to be updated
--> Running transaction check
---> Package perl-Error.noarch 1:0.17010-1.el5 set to be updated
---> Package perl-Git.x86_64 0:1.7.4.1-1.el5 set to be updated
--> Finished Dependency Resolution

Dependencies Resolved

=====
Package Arch Version Repository Size
=====
Installing:
git x86_64 1.7.4.1-1.el5 epel 4.5 M
wget x86_64 1.11.4-2.el5_4.1 base 582 k
Installing for dependencies:
perl-Error noarch 1:0.17010-1.el5 epel 26 k
perl-Git x86_64 1.7.4.1-1.el5 epel 28 k

Transaction Summary
=====
Install 4 Package(s)
Upgrade 0 Package(s)

Total download size: 5.1 M
Downloading Packages:
(1/4): perl-Error-0.17010-1.el5.noarch.rpm | 26 kB 00:00
(2/4): perl-Git-1.7.4.1-1.el5.x86_64.rpm | 28 kB 00:00
(3/4): wget-1.11.4-2.el5_4.1.x86_64.rpm | 582 kB 00:00
(4/4): git-1.7.4.1-1.el5.x86_64.rpm | 4.5 MB 00:01
-----
Total 2.6 MB/s | 5.1 MB 00:02
warning: rpmts_HdrFromFdno: Header V3 DSA signature: NOKEY, key ID 217521f6
epel/gpgkey | 1.7 kB 00:00
Importing GPG key 0x217521F6 "Fedora EPEL <epel@fedoraproject.org>" from /etc/pki/rpm-gpg/RPM-GPG-KEY-EPEL
Running rpm_check_debug
Running Transaction Test
Finished Transaction Test
Transaction Test Succeeded
Running Transaction
Installing : wget 1/4
Installing : perl-Error 2/4
Installing : git 3/4
Installing : perl-Git 4/4

Installed:
git.x86_64 0:1.7.4.1-1.el5 wget.x86_64 0:1.11.4-2.el5_4.1

Dependency Installed:
perl-Error.noarch 1:0.17010-1.el5 perl-Git.x86_64 0:1.7.4.1-1.el5

Complete!
```

```

[cloudera@localhost ~]$ sudo rpm -Uvh http://dl.fedoraproject.org/pub/epel/5/x86\_64/epel-release-5-4.noarch.rpm
Retrieving http://dl.fedoraproject.org/pub/epel/5/x86\_64/epel-release-5-4.noarch.rpm
warning: /var/tmp/rpm-xfer.CPJMIi: Header V3 DSA signature: NOKEY, key ID 217521f6
Preparing... ##### [100%]
 1:epel-release ##### [100%]
[cloudera@localhost ~]$ sudo yum -y install R
Loaded plugins: fastestmirror
Loading mirror speeds from cached hostfile
* base: mirror.symnds.com
* epel: mirrors.einstein.yu.edu
* extras: mirrors.einstein.yu.edu
* updates: mirror.symnds.com
Setting up Install Process
Resolving Dependencies
There are unfinished transactions remaining. You might consider running yum-complete-transaction first to finish them.
The program yum-complete-transaction is found in the yum-utils package.
--> Running transaction check
---> Package R.x86_64 0:2.14.1-1.el5 set to be updated
--> Processing Dependency: libRmath-devel = 2.14.1-1.el5 for package: R
--> Processing Dependency: R-devel = 2.14.1-1.el5 for package: R
--> Running transaction check
---> Package R-devel.x86_64 0:2.14.1-1.el5 set to be updated
--> Processing Dependency: R-core = 2.14.1-1.el5 for package: R-devel
--> Processing Dependency: zlib-devel for package: R-devel
--> Processing Dependency: tk-devel for package: R-devel
--> Processing Dependency: texinfo-tex for package: R-devel
--> Processing Dependency: tetex-latex for package: R-devel
--> Processing Dependency: tcl-devel for package: R-devel
--> Processing Dependency: pcre-devel for package: R-devel
--> Processing Dependency: libX11-devel for package: R-devel
--> Processing Dependency: gcc-gfortran for package: R-devel
--> Processing Dependency: gcc-c++ for package: R-devel
--> Processing Dependency: bzip2-devel for package: R-devel
---> Package libRmath-devel.x86_64 0:2.14.1-1.el5 set to be updated
--> Processing Dependency: libRmath = 2.14.1-1.el5 for package: libRmath-devel
--> Running transaction check
---> Package R-core.x86_64 0:2.14.1-1.el5 set to be updated
--> Processing Dependency: xdg-utils for package: R-core
--> Processing Dependency: cups for package: R-core
--> Processing Dependency: libgfortran.so.1()(64bit) for package: R-core
---> Package bzip2-devel.x86_64 0:1.0.3-6.el5_5 set to be updated
---> Package gcc-c++.x86_64 0:4.1.2-51.el5 set to be updated
--> Processing Dependency: gcc = 4.1.2-51.el5 for package: gcc-c++
--> Processing Dependency: libstdc++-devel = 4.1.2-51.el5 for package: gcc-c++
---> Package gcc-gfortran.x86_64 0:4.1.2-51.el5 set to be updated
--> Processing Dependency: libgmp.so.3()(64bit) for package: gcc-gfortran
---> Package libRmath.x86_64 0:2.14.1-1.el5 set to be updated
---> Package libX11-devel.x86_64 0:1.0.3-11.el5_7.1 set to be updated
--> Processing Dependency: xorg-x11-proto-devel >= 7.1-2 for package: libX11-devel
--> Processing Dependency: libXau-devel for package: libX11-devel
--> Processing Dependency: libXdmcp-devel for package: libX11-devel
---> Package pcre-devel.x86_64 0:6.6-6.el5_6.1 set to be updated
---> Package tcl-devel.x86_64 0:8.4.13-4.el5 set to be updated
---> Package tetex-latex.x86_64 0:3.0-33.13.el5 set to be updated
--> Processing Dependency: tetex-dvips = 3.0 for package: tetex-latex
--> Processing Dependency: tetex = 3.0 for package: tetex-latex
--> Processing Dependency: netpbm-progs for package: tetex-latex
---> Package texinfo-tex.x86_64 0:4.8-14.el5 set to be updated
--> Processing Dependency: texinfo = 4.8-14.el5 for package: texinfo-tex
---> Package tk-devel.x86_64 0:8.4.13-5.el5_1.1 set to be updated
---> Package zlib-devel.x86_64 0:1.2.3-4.el5 set to be updated
--> Running transaction check

```

Pretty impressive for
cut-and-pasting a few
commands, eh?

Test Hadoop with a small job

Download my fork of Jonathan Seidman's sample R code from github

```
$ mkdir hadoop-r
$ cd hadoop-r
$ git init
$ git pull git://github.com/jeffreybreen/hadoop-R.git
```

Grab first 1,000 lines from ASA's 2004 airline data

```
$ curl http://stat-computing.org/dataexpo/2009/2004.csv.bz2 | bzip2 \
  | head -1000 > 2004-1000.csv
```

Make some directories in HDFS and load the data file

```
$ hadoop fs -mkdir /user/cloudera
$ hadoop fs -mkdir asa-airline
$ hadoop fs -mkdir asa-airline/data
$ hadoop fs -mkdir asa-airline/out
$ hadoop fs -put 2004-1000.csv asa-airline/data/
```

Run Jonathan's sample streaming job

```
$ cd airline/src/deptdelay_by_month/R/streaming
$ hadoop jar /usr/lib/hadoop/contrib/streaming/hadoop-streaming-*.jar \
  -input asa-airline/data -output asa-airline/out/dept-delay-month \
  -mapper map.R -reducer reduce.R -file map.R -file reduce.R
```

```
[cloudera@localhost hadoop-r]$ head -2 2004-1000.csv
Year,Month,DayofMonth,DayOfWeek,DepTime,CRSDepTime,ArrTime,CRSArrTime,UniqueCarrier,FlightNum,TailNum,ActualElapsedTime,CRSElapsedTime,AirTime,ArrDelay,DepDelay,Origin,Dest,Distance,TaxiIn,TaxiOut,Cancelled,CancellationCode,Diverted,CarrierDelay,WeatherDelay,NASDelay,SecurityDelay,LateAircraftDelay
2004,1,12,1,623,630,901,915,UA,462,N805UA,98,105,80,-14,-7,ORD,CLT,599,7,11,0,,0,0,0,0,0,0
[cloudera@localhost hadoop-r]$ tail -2 2004-1000.csv
2004,1,25,7,857,900,1441,1446,UA,484,N457UA,224,226,208,-5,-3,PDX,ORD,1739,5,11,0,,0,0,0,0,0,0
2004,1,26,1,903,900,1524,1444,UA,484,N554UA,261,224,200,40,3,PDX,ORD,1739,25,36,0,,0,0,0,40,0,0
```

```
[cloudera@localhost hadoop-r]$ cd airline/src/deptdelay_by_month/R/streaming
```

```
[cloudera@localhost streaming]$ hadoop jar /usr/lib/hadoop/contrib/streaming/hadoop-streaming-*.jar \
> -input asa-airline/data -output asa-airline/out/dept-delay-month \
> -mapper map.R -reducer reduce.R -file map.R -file reduce.R
packageJobJar: [map.R, reduce.R, /var/lib/hadoop-0.20/cache/cloudera/hadoop-unjar4442605735512091493/] [] /tmp/
streamjob2138397329652275361.jar tmpDir=null
12/03/06 15:28:15 WARN snappy.LoadSnappy: Snappy native library is available
12/03/06 15:28:15 INFO util.NativeCodeLoader: Loaded the native-hadoop library
12/03/06 15:28:15 INFO snappy.LoadSnappy: Snappy native library loaded
12/03/06 15:28:15 INFO mapred.FileInputFormat: Total input paths to process : 1
12/03/06 15:28:17 INFO streaming.StreamJob: getLocalDirs(): [/var/lib/hadoop-0.20/cache/cloudera/mapred/local]
12/03/06 15:28:17 INFO streaming.StreamJob: Running job: job_201203061110_0001
12/03/06 15:28:17 INFO streaming.StreamJob: To kill this job, run:
12/03/06 15:28:17 INFO streaming.StreamJob: /usr/lib/hadoop-0.20/bin/hadoop job -Dmapred.job.tracker=0.0.0.0:8021 -kill
job_201203061110_0001
12/03/06 15:28:17 INFO streaming.StreamJob: Tracking URL: http://0.0.0.0:50030/jobdetails.jsp?
jobid=job\_201203061110\_0001
12/03/06 15:28:18 INFO streaming.StreamJob: map 0% reduce 0%
12/03/06 15:28:37 INFO streaming.StreamJob: map 100% reduce 0%
12/03/06 15:29:15 INFO streaming.StreamJob: map 100% reduce 100%
12/03/06 15:29:18 INFO streaming.StreamJob: Job complete: job_201203061110_0001
12/03/06 15:29:18 INFO streaming.StreamJob: Output: asa-airline/out/dept-delay-month
```

```
[cloudera@localhost streaming]$ hadoop fs -ls asa-airline/out/dept-delay-month
```

```
Found 3 items
```

```
-rw-r--r-- 1 cloudera supergroup 0 2012-03-06 15:29 /user/cloudera/asa-airline/out/dept-delay-month/_SUCCESS
drwxr-xr-x - cloudera supergroup 0 2012-03-06 15:28 /user/cloudera/asa-airline/out/dept-delay-month/_logs
-rw-r--r-- 1 cloudera supergroup 33 2012-03-06 15:29 /user/cloudera/asa-airline/out/dept-delay-month/
part-00000
```

```
[cloudera@localhost streaming]$ hadoop fs -cat asa-airline/out/dept-delay-month/part-00000
```

```
2004 1 973 UA 11.55293
```


Install RHadoop's **rmr** package

- RHadoop is an open source project sponsored by Revolution Analytics and is one of several available to make it easier to work with R and Hadoop
 - The **rmr** package contains all the mapreduce-related functions, including generating Hadoop streaming jobs and basic data exchange with HDFS
- First install prerequisite packages (run R as root to install system-wide)

```
$ sudo R
> install.packages( c('RJSONIO', 'itertools', 'digest'),
  repos='http://cran.revolutionanalytics.com')
```

- Download the latest stable release (1.2) from github

```
$ wget --no-check-certificate https://github.com/downloads/RevolutionAnalytics/RHadoop/rmr_1.2.tar.gz
```

- Install the package from the tar file

```
$ sudo R CMD INSTALL rmr_1.2.tar.gz
```

- Test that it loads

```
$ R
> library(rmr)
Loading required package: RJSONIO
Loading required package: itertools
Loading required package: iterators
Loading required package: digest
```

Test **rmr** with the airline example

- Runs same analysis as streaming example, but using **rmr**'s abstractions

```
$ cd
$ cd hadoop-r/airline/src/deptdelay_by_month/R/rmr/
$ export HADOOP_HOME=/usr/lib/hadoop
$ R
[...]  
> source('deptdelay-rmr12.R')
```

- It will fail because our HDFS input paths don't match, but it did load all the functions so we can easily kick off the job by hand:

```
> df = from.dfs(deptdelay("asa-airline/data", "asa-airline/out/deptdelay-  
month-rmr"), to.data.frame=T)  
[...]  
> colnames(df) = c('year', 'month', 'count', 'airline', 'mean.delay')  
> df
```

	year	month	count	airline	mean.delay
rmr.key	2004	1	973	UA	11.5529290853032

```
> df = from.dfs(deptdelay("asa-airline/data", "asa-airline/out/deptdelay-month-rmr"),
to.data.frame=T)
```

```
packageJobJar: [/tmp/RtmpZAckHy/rhstr.map4da957c5e126, /tmp/RtmpZAckHy/
rhstr.reduce4da938d5ffcb, /tmp/RtmpZAckHy/rmr-local-env, /tmp/RtmpZAckHy/rmr-global-env, /
var/lib/hadoop-0.20/cache/cloudera/hadoop-unjar674649393612449255/] [] /tmp/
streamjob8188313657687081754.jar tmpDir=null
```

```
12/03/06 16:28:57 WARN snappy.LoadSnappy: Snappy native library is available
```

```
12/03/06 16:28:57 INFO util.NativeCodeLoader: Loaded the native-hadoop library
```

```
12/03/06 16:28:57 INFO snappy.LoadSnappy: Snappy native library loaded
```

```
12/03/06 16:28:57 INFO mapred.FileInputFormat: Total input paths to process : 1
```

```
12/03/06 16:28:58 INFO streaming.StreamJob: getLocalDirs(): [/var/lib/hadoop-0.20/cache/
cloudera/mapred/local]
```

```
12/03/06 16:28:58 INFO streaming.StreamJob: Running job: job_201203061110_0003
```

```
12/03/06 16:28:58 INFO streaming.StreamJob: To kill this job, run:
```

```
12/03/06 16:28:58 INFO streaming.StreamJob: /usr/lib/hadoop/bin/hadoop job -
Dmapred.job.tracker=0.0.0.0:8021 -kill job_201203061110_0003
```

```
12/03/06 16:28:58 INFO streaming.StreamJob: Tracking URL: http://0.0.0.0:50030/
jobdetails.jsp?jobid=job\_201203061110\_0003
```

```
12/03/06 16:28:59 INFO streaming.StreamJob: map 0% reduce 0%
```

```
12/03/06 16:29:21 INFO streaming.StreamJob: map 100% reduce 0%
```

```
12/03/06 16:29:46 INFO streaming.StreamJob: map 100% reduce 100%
```

```
12/03/06 16:29:55 INFO streaming.StreamJob: Job complete: job_201203061110_0003
```

```
12/03/06 16:29:55 INFO streaming.StreamJob: Output: asa-airline/out/deptdelay-month-rmr
```

```
>
```

```
> colnames(df) = c('year', 'month', 'count', 'airline', 'mean.delay')
```

```
> df
```

	year	month	count	airline	mean.delay
rmr.key	2004	1	973	UA	11.5529290853032

Extra Credit: Install RStudio

- Current download link and instructions at <http://rstudio.org/download/server>

```
$ wget http://download2.rstudio.org/rstudio-server-0.95.262-x86_64.rpm
```

```
$ sudo rpm -Uvh rstudio-server-0.95.262-x86_64.rpm
```

- Find IP address with ifconfig

```
$ ifconfig
```

- Access from browser via port 8787

- e.g., <http://192.168.1.140:8787/>

```
[cloudera@localhost ~]$ wget http://download2.rstudio.org/rstudio-server-0.95.262-x86_64.rpm
--2012-03-06 12:14:24-- http://download2.rstudio.org/rstudio-server-0.95.262-x86_64.rpm
Resolving download2.rstudio.org... 216.137.39.181, 216.137.39.217, 216.137.39.222, ...
Connecting to download2.rstudio.org|216.137.39.181|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 15748959 (15M) [application/x-redhat-package-manager]
Saving to: `rstudio-server-0.95.262-x86_64.rpm'

100%[=====>] 15,748,959 1.83M/s in 7.2s

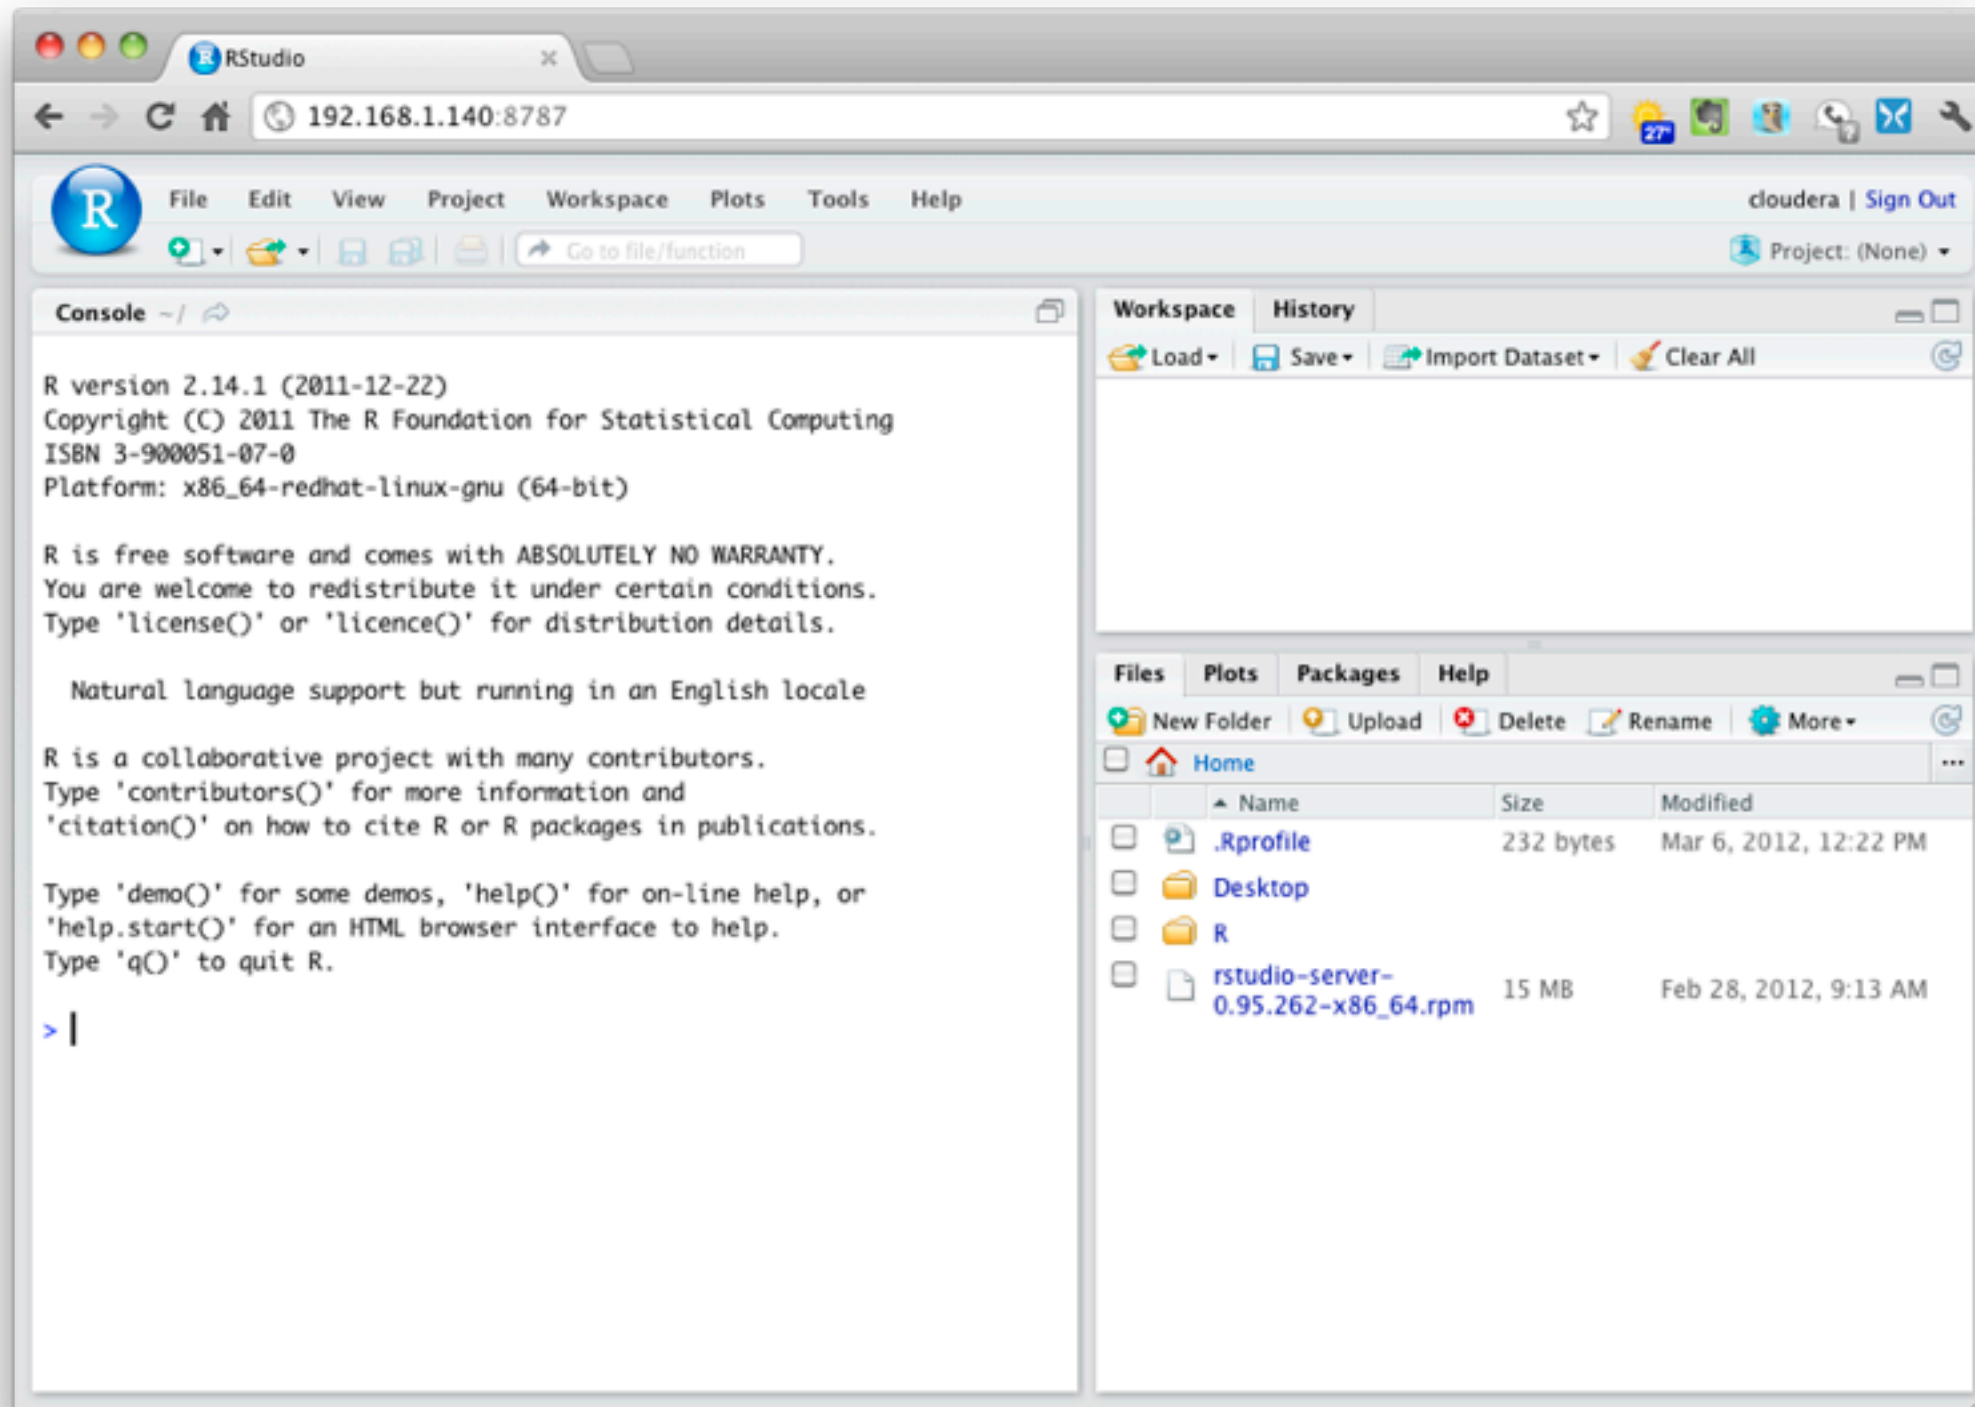
2012-03-06 12:14:31 (2.09 MB/s) - `rstudio-server-0.95.262-x86_64.rpm' saved [15748959/15748959]

[cloudera@localhost ~]$ sudo rpm -Uvh rstudio-server-0.95.262-x86_64.rpm
Preparing... ##### [100%]
 1:rstudio-server ##### [100%]
rsession: no process killed
Starting rstudio-server: [ OK ]
[cloudera@localhost ~]$ ifconfig

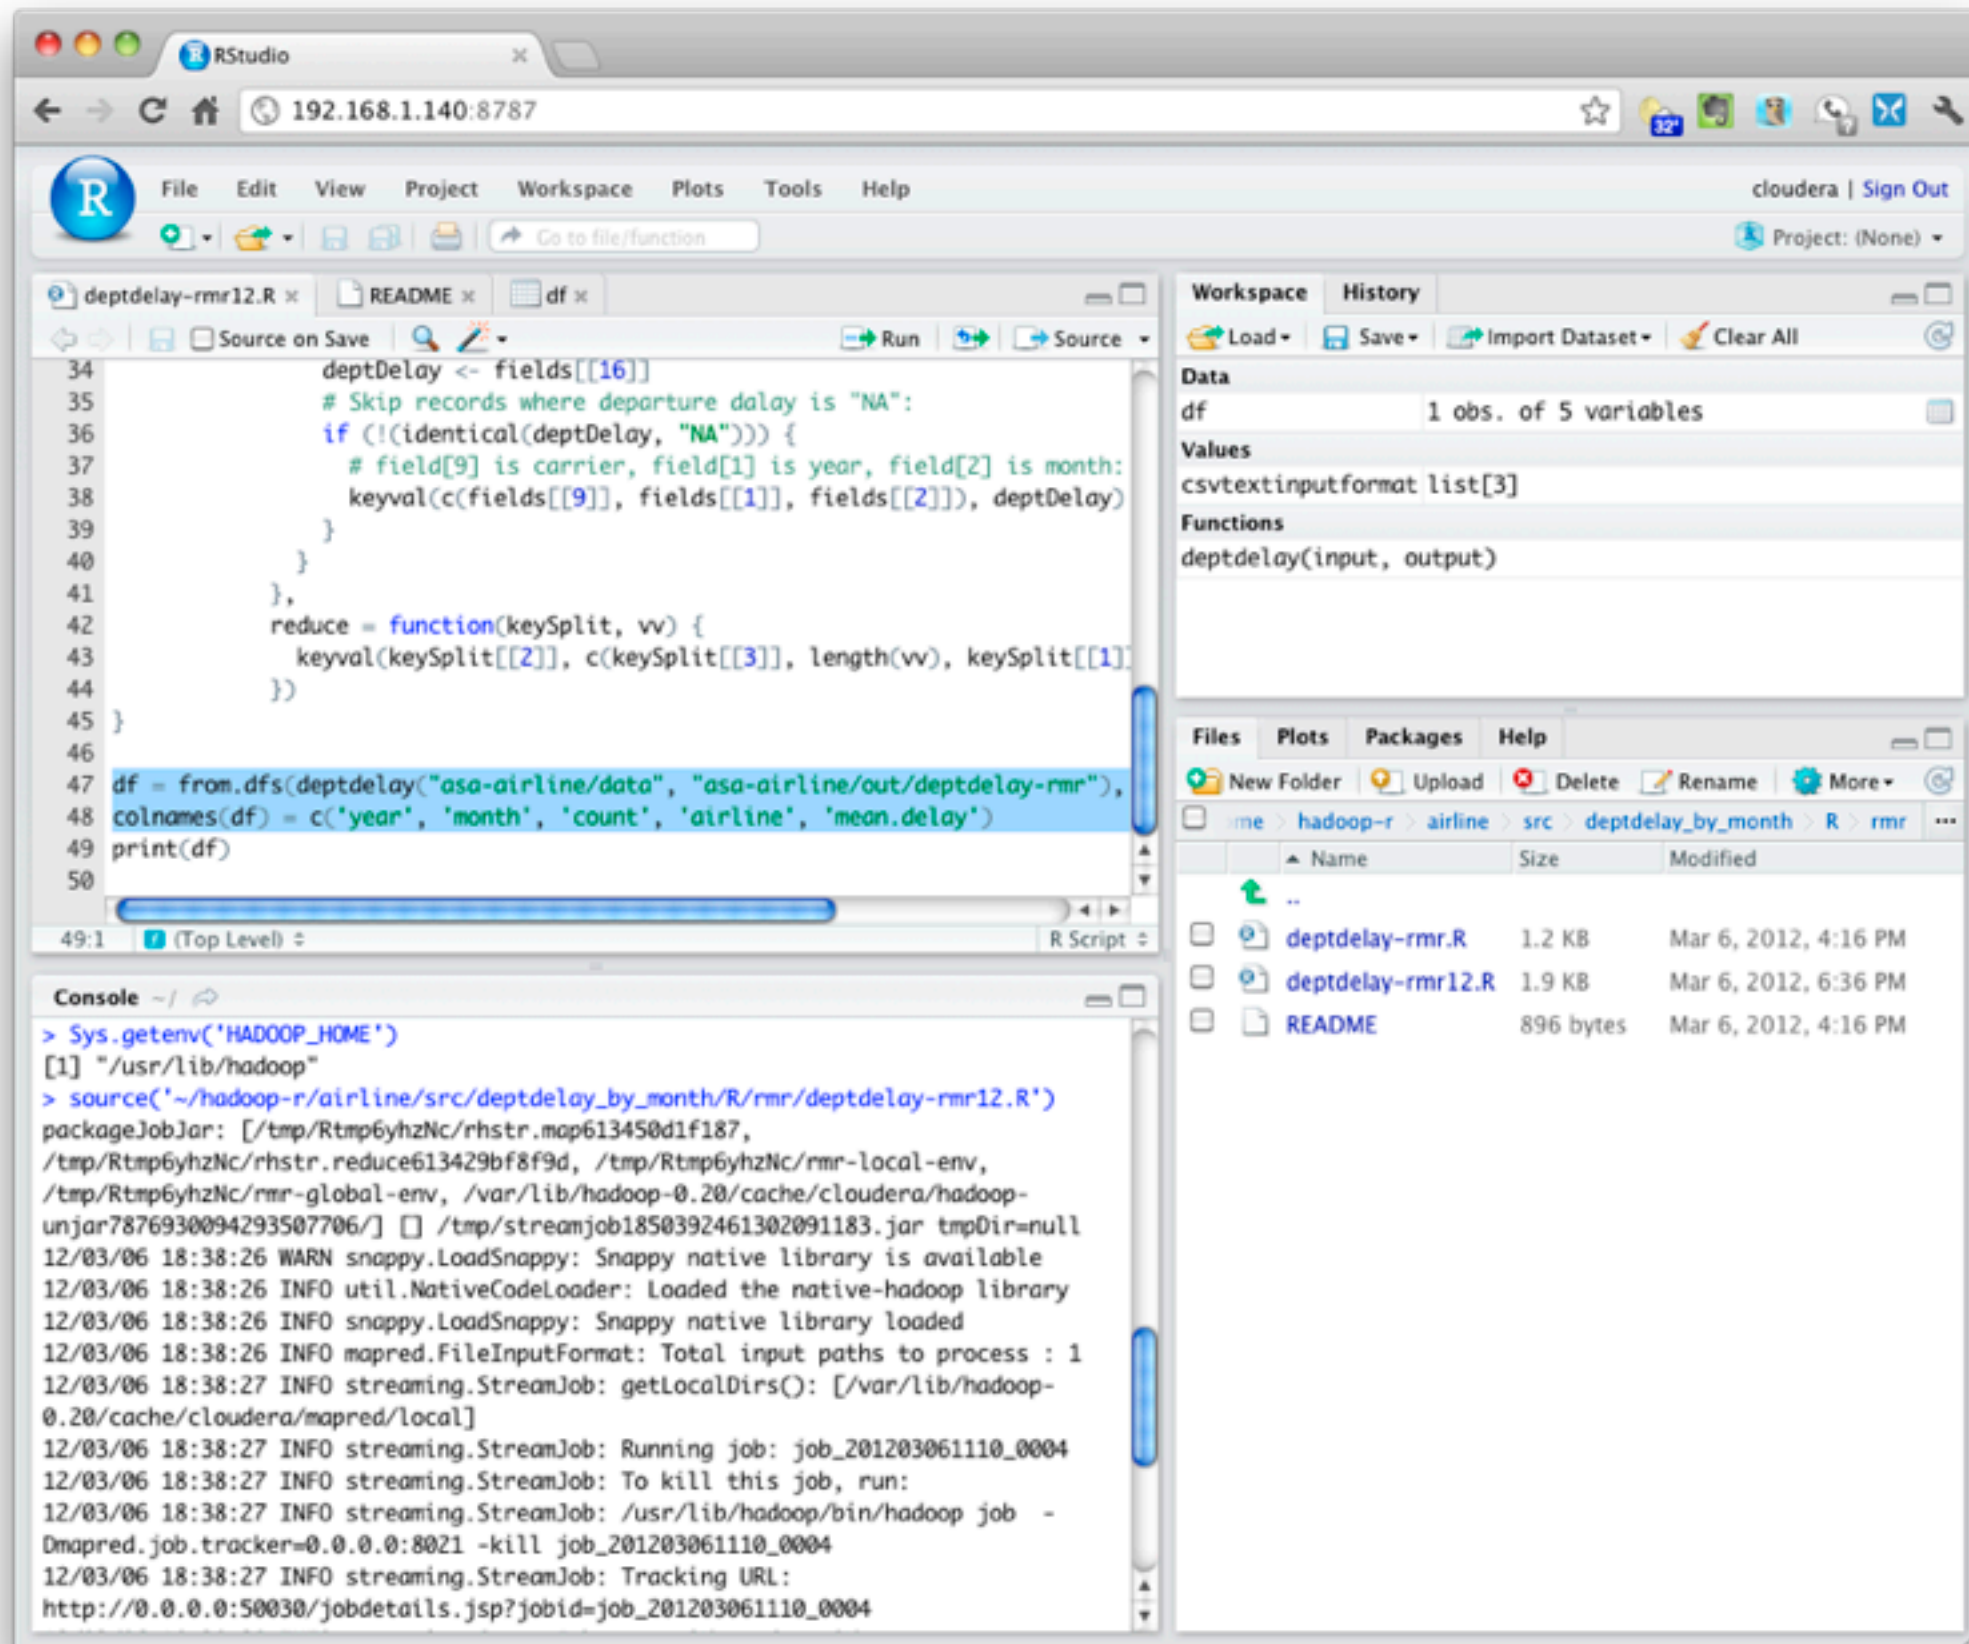
eth0      Link encap:Ethernet  HWaddr 00:0C:29:4B:77:1D
          inet addr:192.168.1.140  Bcast:192.168.1.255  Mask:255.255.255.0
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:75039 errors:0 dropped:0 overruns:0 frame:0
          TX packets:36742 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:104953280 (100.0 MiB)  TX bytes:3061577 (2.9 MiB)
          Interrupt:59 Base address:0x2000

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          UP LOOPBACK RUNNING  MTU:16436  Metric:1
          RX packets:78954 errors:0 dropped:0 overruns:0 frame:0
          TX packets:78954 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:14608044 (13.9 MiB)  TX bytes:14608044 (13.9 MiB)
```

RStudio Success



RStudio + **rmr** works too



Next up:
Running R & RStudio
on Amazon EC2